

The Linux graphics stack, Optimus and the Nouveau driver

Cooperative rendering across GPUs on Linux

Martin Peres

Nouveau developer
PhD student at LaBRI
X.Org Foundation board member

September 24, 2014

Summary

- 1 Introduction to the Linux graphics stack
 - General overview
 - Kernel space
 - User space
- 2 Optimus
- 3 Kernel
- 4 Userspace
- 5 Tools
- 6 Community

General overview of the Linux Graphics stack

The graphics stack before 2005

- The X-Server provided everything:
 - Modesetting (CRTC & plane management);
 - 2D/3D acceleration;
 - Video rendering acceleration;
 - Input management.
- The X-Server talked to the GPU directly, as root.

The current graphics stack

- The X-Server got split into more than 200 components:
 - Privileged operations moved to the kernel;
 - 2D drivers got put into different shared objects;
 - 3D acceleration got put in mesa;
 - The list is too long ;)

by Shmuel Csaba Otto Traian; CC-BY-SA 4.0 intl; created 2013-08-24; last updated 2014-03-25

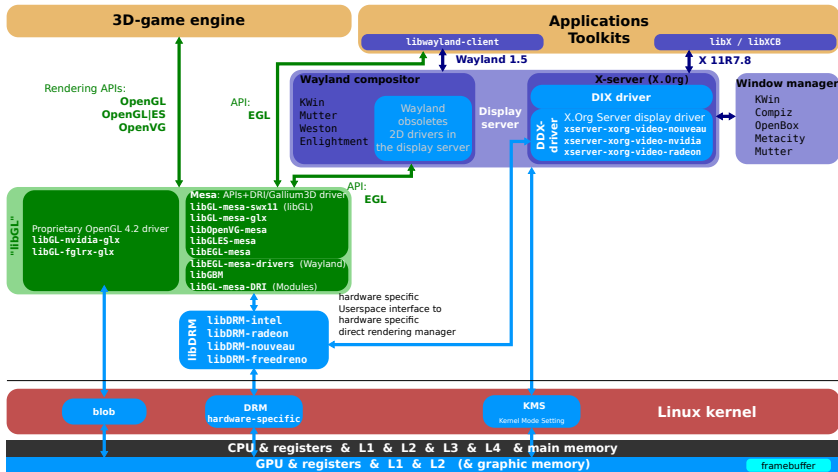


Figure: General overview of the Linux graphics stack

The kernel space

Direct Rendering Manager (DRM) : The common code

- This common code provides:
 - Kernel ModeSetting (KMS): CRTC & plane management;
 - Video memory management via GEM (with a TTM backend?);
 - Nodes with different capabilities (master or render nodes).

DRM open source drivers

- i810/i915: Intel;
- nouveau: NVIDIA;
- radeon: AMD/ATI;
- vmwgfx: VMware;
- many SoC GPUs (armada, exynos, msm, omap, tegra, ...).

Architecture of the X-Server

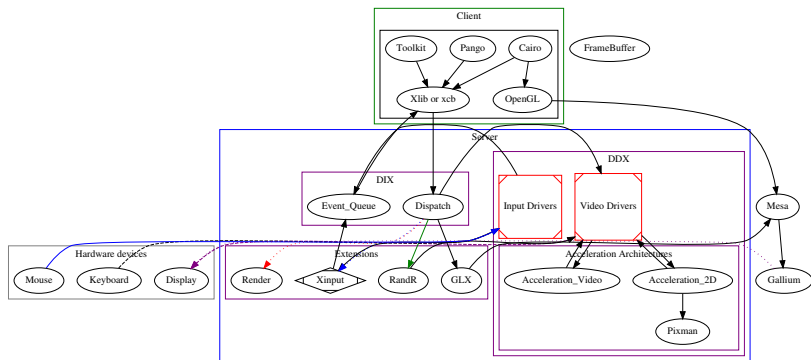


Figure: General overview of the X-Server's internal architecture

Architecture of Mesa




Figure: General overview of Mesa's internal architecture

Summary

1 Introduction to the Linux graphics stack

2 **Optimus**

- Introduction
- Turning the dGPU on/off
- Driving the right outputs

3 Kernel

4 Userspace

5 Tools

6 Community

Great performance, great battery-life

Optimus

- Laptops can be equipped with two GPUs;
- The Intel IGP is great for battery-life;
- NVIDIA's discrete GPU (dGPU) is great for performance;
- Dynamic switch between the 2: get the best of both worlds!

Challenges

- When/How the dGPU should be turned on/off?
- Who drives the outputs?
- How to copy buffers from a driver to another?
- How should we handle the HDMI "sound card"?

Turning the dGPU on/off

How

- Optimus laptops have ACPI functions to do that;
- Two ways of calling them:
 - bbswitch: Old kernel module for manual management;
 - vgaswitcheroo: Manual or automatic state management.

When: The case of vgaswitcheroo

- Turn off the dGPU when it has been idle for 5 seconds;
- Idle?:
 - no graphics context allocated;
 - no output is being used;
 - no sound interface used (not done);
 - no call to the drm driver has been made;

Handling the outputs : Hardware multiplexer

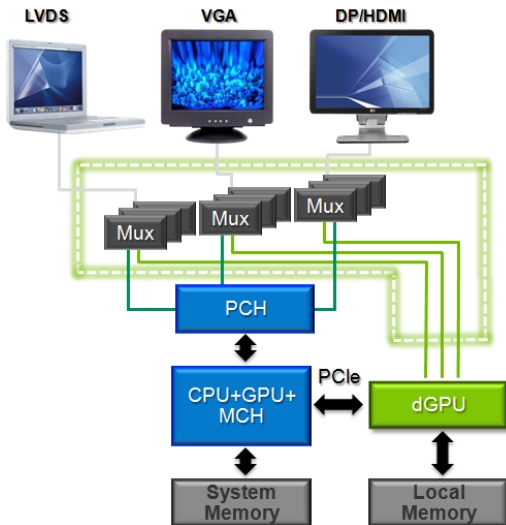


Figure: Switchable graphics

Handling the outputs : Software multiplexer

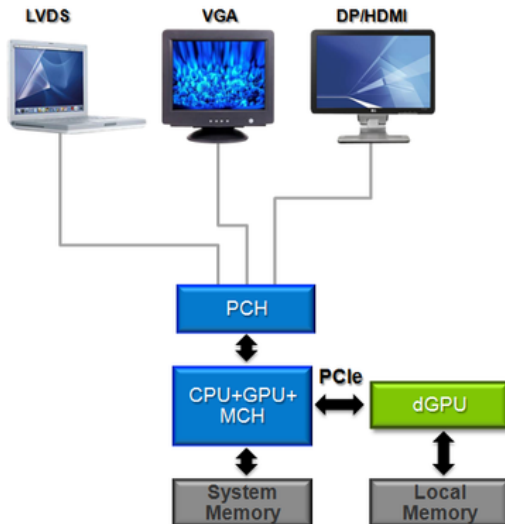


Figure: The “real” optimus architecture

Switching from one GPU to another : How windows does it




Figure: The global hardware/software infrastructure

Sharing buffers across drivers

Challenges

- The memory representation for buffers is different from hardware to hardware:
 - pitch: number of pixels per row;
 - tiling: technique that increases the spatial locality.
- Synchronising rendering across drivers.

DMA-Buf

The memory representation for buffers is different from hardware to hardware:

- pitch: number of pixels per row;
- tiling: technique that increases the spatial locality.

Summary

- 1 Introduction to the Linux graphics stack
- 2 Optimus
- 3 Kernel**
 - Optimus/prime
 - Power Management
- 4 Userspace
- 5 Tools
- 6 Community

Prime

Prime

Prime is the name for all the open source technologies that make hybrid graphics possible:

- vgaswitcheroo: switching graphics ;
- running the nouveau ddx;

List of requirements

- running nouveau/radeon drm;
- running the nouveau ddx;